

Navigation

April 26, 2020

1 Navigation

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

1.0.1 1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

```
In [1]: !pip -q install ./python
```

```
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 3.0.
```

The environment is already saved in the Workspace and can be accessed at the file path provided below. Please run the next code cell without making any changes.

```
In [2]: import numpy as np
import torch
import matplotlib.pyplot as plt

from collections import deque
from unityagents import UnityEnvironment
from agent import Agent

# please do not modify the line below
env = UnityEnvironment(file_name="/data/Banana_Linux_NoVis/Banana.x86_64")
```

```
INFO:unityagents:
'Academy' started successfully!
Unity Academy name: Academy
  Number of Brains: 1
  Number of External Brains : 1
  Lesson number : 0
  Reset Parameters :
```

```

Unity brain name: BananaBrain
  Number of Visual Observations (per agent): 0
  Vector Observation space type: continuous
  Vector Observation space size (per agent): 37
  Number of stacked Vector Observation: 1
  Vector Action space type: discrete
  Vector Action space size (per agent): 4
  Vector Action descriptions: , , ,

```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```

In [3]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]

```

1.0.2 2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

```

In [4]: # reset the environment
        env_info = env.reset(train_mode=True)[brain_name]

        # number of agents in the environment
        print('Number of agents:', len(env_info.agents))

        # number of actions
        action_size = brain.vector_action_space_size
        print('Number of actions:', action_size)

        # examine the state space
        state = env_info.vector_observations[0]
        print('States look like:', state)
        state_size = len(state)
        print('States have length:', state_size)

```

```

Number of agents: 1
Number of actions: 4
States look like: [1.          0.          0.          0.          0.84408134 0.
 0.          1.          0.          0.0748472 0.          1.
 0.          0.          0.25755   1.          0.          0.
 0.          0.74177343 0.          1.          0.          0.
 0.25854847 0.          0.          1.          0.          0.09355672
 0.          1.          0.          0.          0.31969345 0.
 0.          ]
States have length: 37

```

1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that **in this coding environment, you will not be able to watch the agent while it is training**, and you should set `train_mode=True` to restart the environment.

```
In [5]: env_info = env.reset(train_mode=False)[brain_name] # reset the environment
        state = env_info.vector_observations[0]             # get the current state
        score = 0                                           # initialize the score
        while True:
            action = np.random.randint(action_size)         # select an action
            env_info = env.step(action)[brain_name]          # send the action to the environment
            next_state = env_info.vector_observations[0]      # get the next state
            reward = env_info.rewards[0]                     # get the reward
            done = env_info.local_done[0]                    # see if episode has finished
            score += reward                                   # update the score
            state = next_state                                # roll over the state to next time step
            if done:                                          # exit loop if episode finished
                break

        print("Score: {}".format(score))
```

Score: 0.0

1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**: - When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on *Jupyter* in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agent while it is training. However, *after training the agent*, you can download the saved model weights to watch the agent on your own machine!

```
In [6]: def dq_n(n_episodes=2000, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=0.995,
               train_mode=True, ckpt_path='checkpoint.pth'):
        """Deep Q-Learning.

        Params
        =====
        n_episodes (int): maximum number of training episodes
        max_t (int): maximum number of timesteps per episode
```

```

    eps_start (float): starting value of epsilon, for epsilon-greedy action selection
    eps_end (float): minimum value of epsilon
    eps_decay (float): multiplicative factor (per episode) for decreasing epsilon
    train_mode (bool): run training mode if `True`
"""
scores = [] # list containing scores from each episode
scores_window = deque(maxlen=100) # last 100 scores
eps = eps_start # initialize epsilon

for i_episode in range(1, n_episodes+1):
    env_info = env.reset(train_mode=train_mode)[brain_name] # reset environment
    state = env_info.vector_observations[0] # get current state
    score = 0
    for t in range(max_t):
        action = agent.act(state, eps) # select an action
        env_info = env.step(action)[brain_name] # send action to environment
        next_state = env_info.vector_observations[0] # get next state
        reward = env_info.rewards[0] # get reward
        done = env_info.local_done[0] # see if episode has finished
        agent.step(state, action, reward, next_state, done) # learning step
        state = next_state
        score += reward
        if done:
            break

    scores_window.append(score) # save most recent score to window
    scores.append(score) # save most recent score to total
    eps = max(eps_end, eps_decay*eps) # decrease epsilon

    print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
    if i_episode % 100 == 0:
        print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
    if np.mean(scores_window) >= 13.0:
        print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
        if train_mode: torch.save(agent.qnetwork_local.state_dict(), ckpt_path)
        break
return scores

```

DQN

```


In [9]: agent = Agent(state_size=state_size, action_size=action_size, seed=0)
        scores = dqn(n_episodes=500, eps_decay=0.98, ckpt_path='v1_checkpoint.pth')

        # plot the scores
        fig = plt.figure()
        ax = fig.add_subplot(111)
        plt.plot(np.arange(len(scores)), scores, label='DQN')
        plt.ylabel('Score')

```

```
plt.xlabel('Episode #')
plt.legend(loc='upper left')
plt.show()
```

Episode 100	Average Score: 0.18
Episode 200	Average Score: 0.68
Episode 300	Average Score: 0.59
Episode 400	Average Score: 0.53
Episode 500	Average Score: 0.60



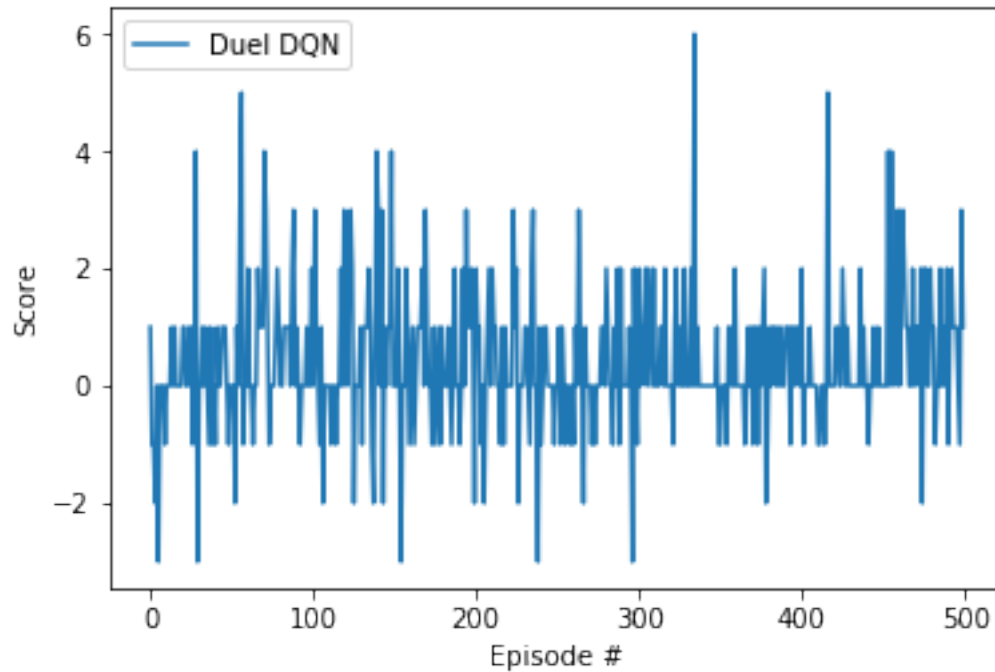
Dueling DQN

```
In [7]: agent = Agent(state_size=state_size, action_size=action_size, seed=0, duel=True)
        scores = dqn(n_episodes=500, eps_decay=0.98, ckpt_path='v2_checkpoint.pth')
```

```
# plot the scores
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores, label='Duel DQN')
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.legend(loc='upper left')
plt.show()
```

Episode 100	Average Score: 0.37
Episode 200	Average Score: 0.51

Episode 300	Average Score: 0.23
Episode 400	Average Score: 0.44
Episode 500	Average Score: 0.60




Double DQN

```
In [7]: agent = Agent(state_size=state_size, action_size=action_size, seed=0, double=True)
        scores = dqn(eps_decay=0.98, ckpt_path='v3_checkpoint.pth')
```

```
# plot the scores
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores, label='DDQN')
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.legend(loc='upper left')
plt.show()
```

Episode 100	Average Score: 1.75
Episode 200	Average Score: 6.07
Episode 300	Average Score: 9.88
Episode 396	Average Score: 13.07
Environment solved in 296 episodes!	Average Score: 13.07



DDQN + Prioritized Experience Replay

```
In [ ]: agent = Agent(state_size=state_size, action_size=action_size, seed=0, double=True, prioritized=True)
scores = dqn(n_episodes=500, eps_decay=0.98, ckpt_path='v4_checkpoint.pth')
```

```
# plot the scores
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores, label='DDQN + PER')
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.legend(loc='upper left')
plt.show()
```

```
Episode 100      Average Score: 0.77
Episode 159      Average Score: 2.41
```


Dueling DDQN

```
In [6]: agent = Agent(state_size=state_size, action_size=action_size, seed=0, double=True, dueling=True)
scores = dqn(eps_decay=0.98, ckpt_path='v5_checkpoint.pth')
```

```
# plot the scores
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores, label='Duel DDQN')
```

```
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.legend(loc='upper left')
plt.show()
```

```
Episode 100      Average Score: 0.23
Episode 200      Average Score: 3.45
Episode 300      Average Score: 6.54
Episode 400      Average Score: 11.34
Episode 453      Average Score: 13.01
Environment solved in 353 episodes!      Average Score: 13.01
```



```
In [ ]: env.close()
```